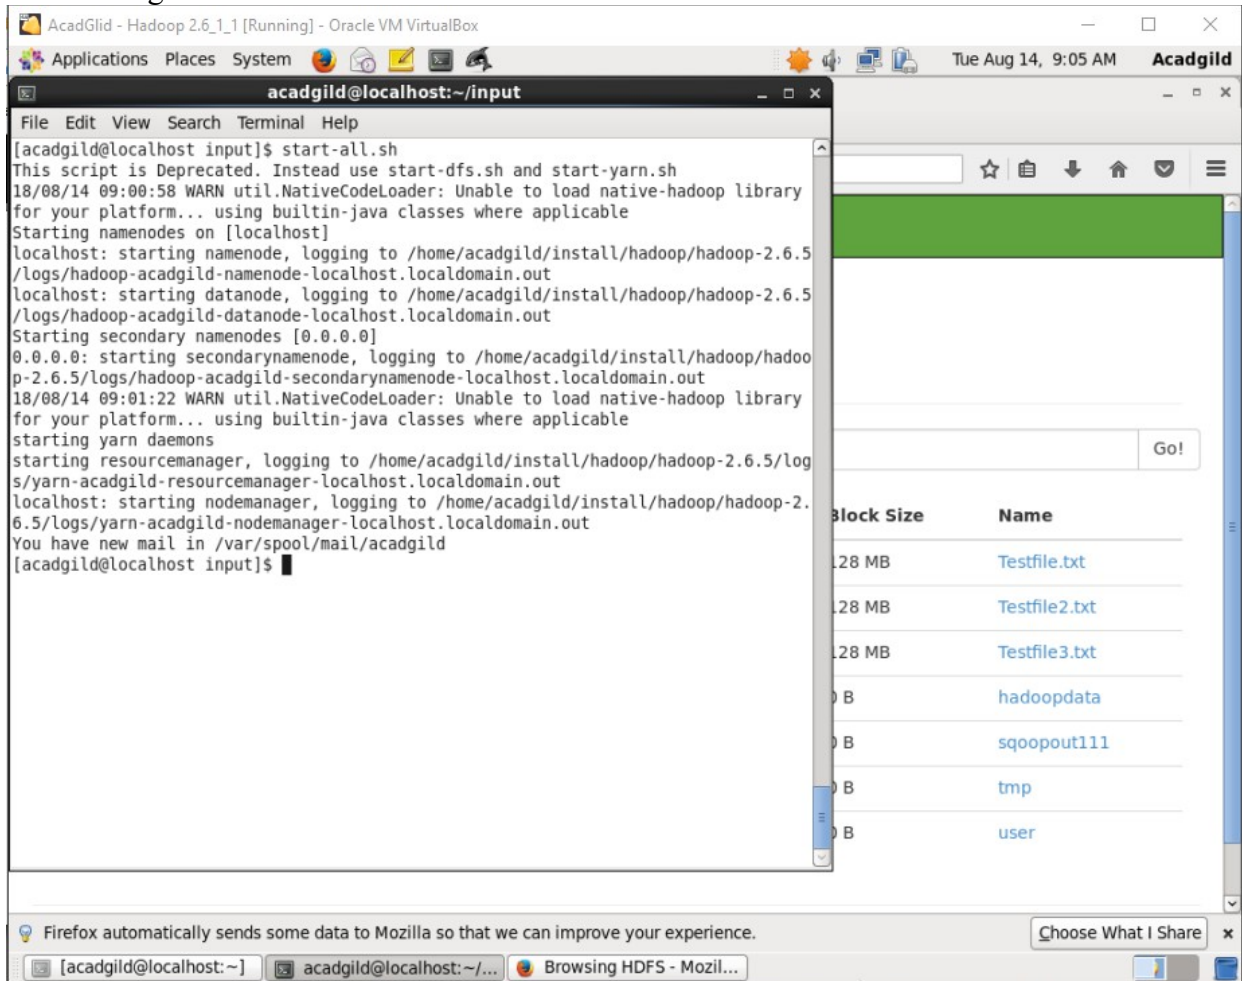


Session 1 – Assignment 1

1. Start-all.sh

Script will start Namenode, Secondary Namenode, Datanode, yarn daemons, resource manager



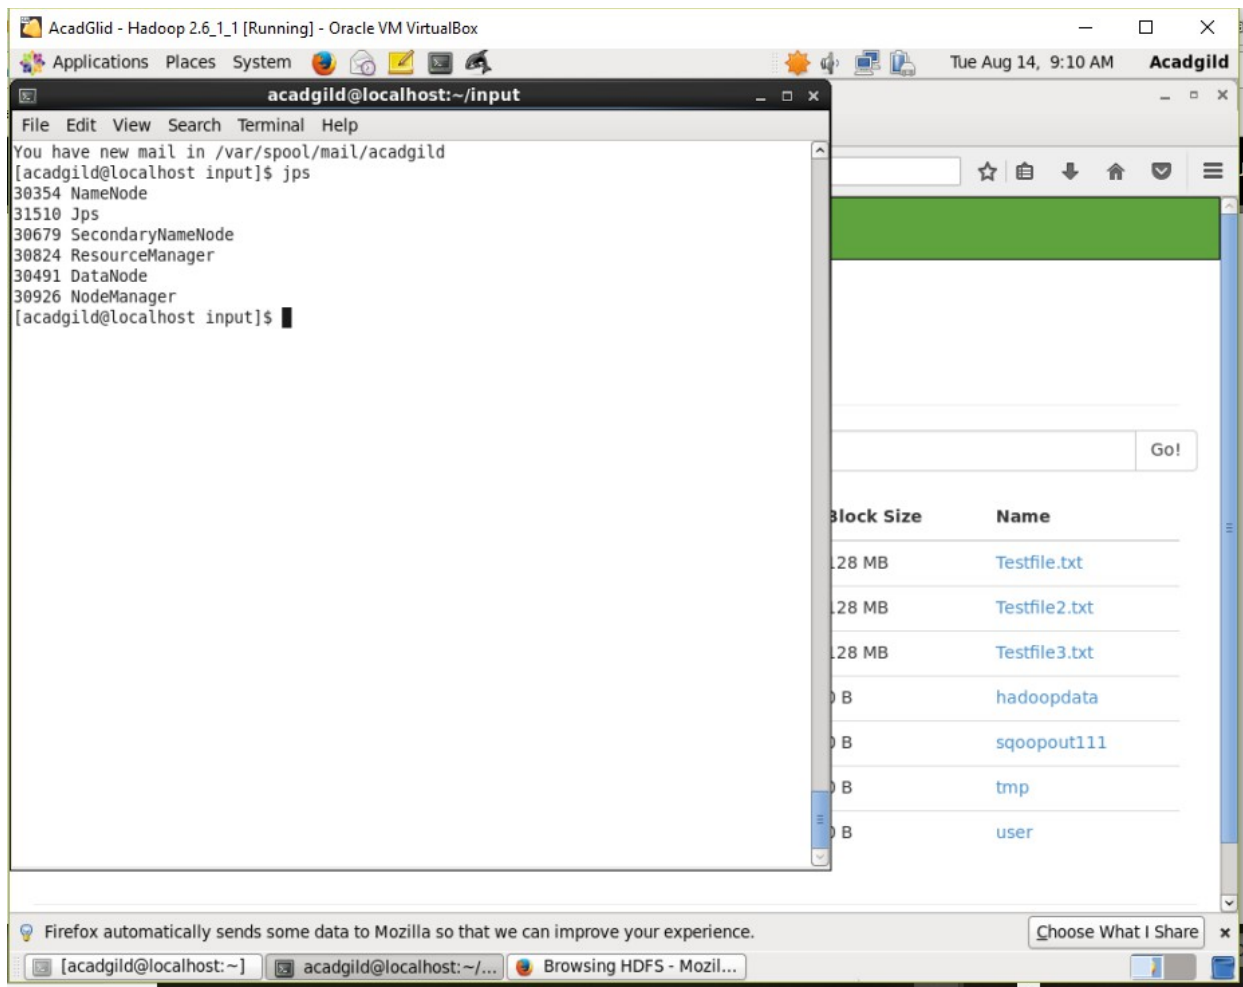
The screenshot shows a terminal window titled 'acadgild@localhost:~/input' with the following output:

```
[acadgild@localhost input]$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
18/08/14 09:00:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5
/logs/hadoop-acadgild-namenode-localhost.localdomain.out
localhost: starting datanode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5
/logs/hadoop-acadgild-datanode-localhost.localdomain.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/acadgild/install/hadoop/hadoo
p-2.6.5/logs/hadoop-acadgild-secondarynamenode-localhost.localdomain.out
18/08/14 09:01:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/log
s/yarn-acadgild-resourcemanager-localhost.localdomain.out
localhost: starting nodemanager, logging to /home/acadgild/install/hadoop/hadoop-2.
6.5/logs/yarn-acadgild-nodemanager-localhost.localdomain.out
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost input]$
```

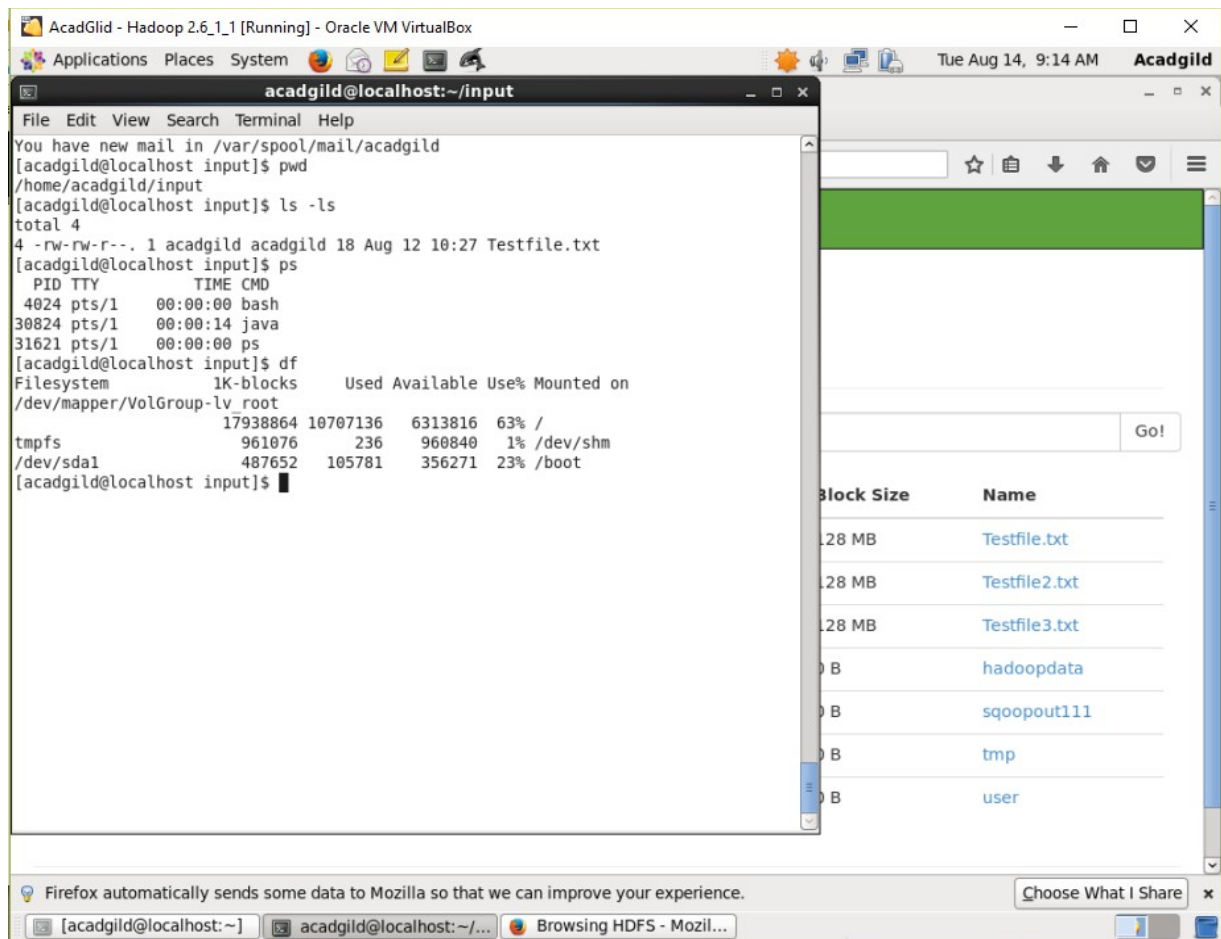
To the right of the terminal, a file browser window shows the contents of the /var/spool/mail/acadgild directory. It displays a table with two columns: 'Block Size' and 'Name'.

Block Size	Name
28 MB	Testfile.txt
28 MB	Testfile2.txt
28 MB	Testfile3.txt
0 B	hadoopdata
0 B	sqoopout111
0 B	tmp
0 B	user

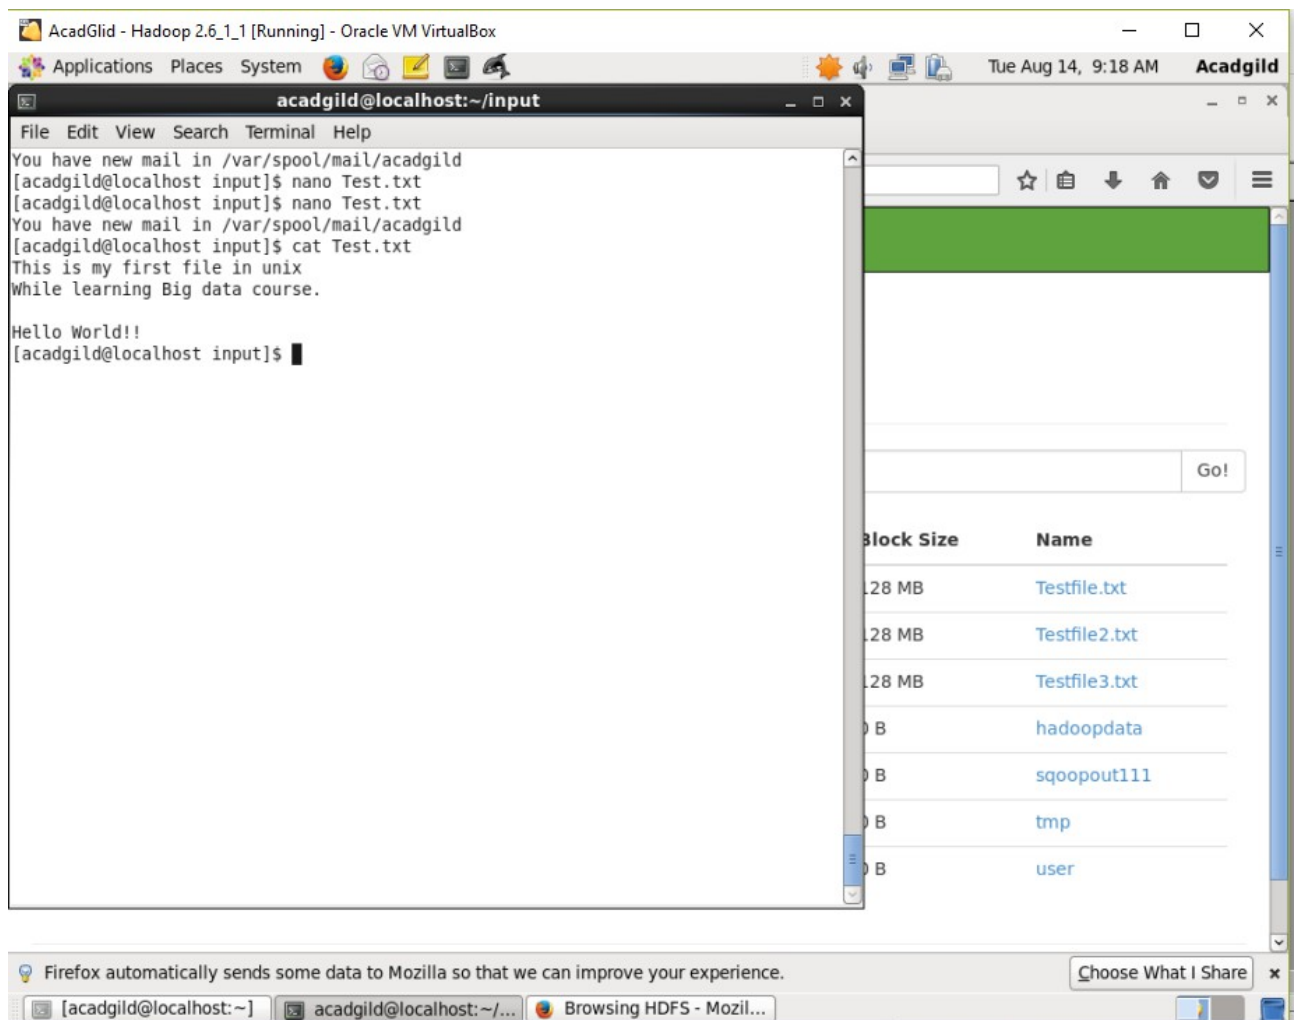
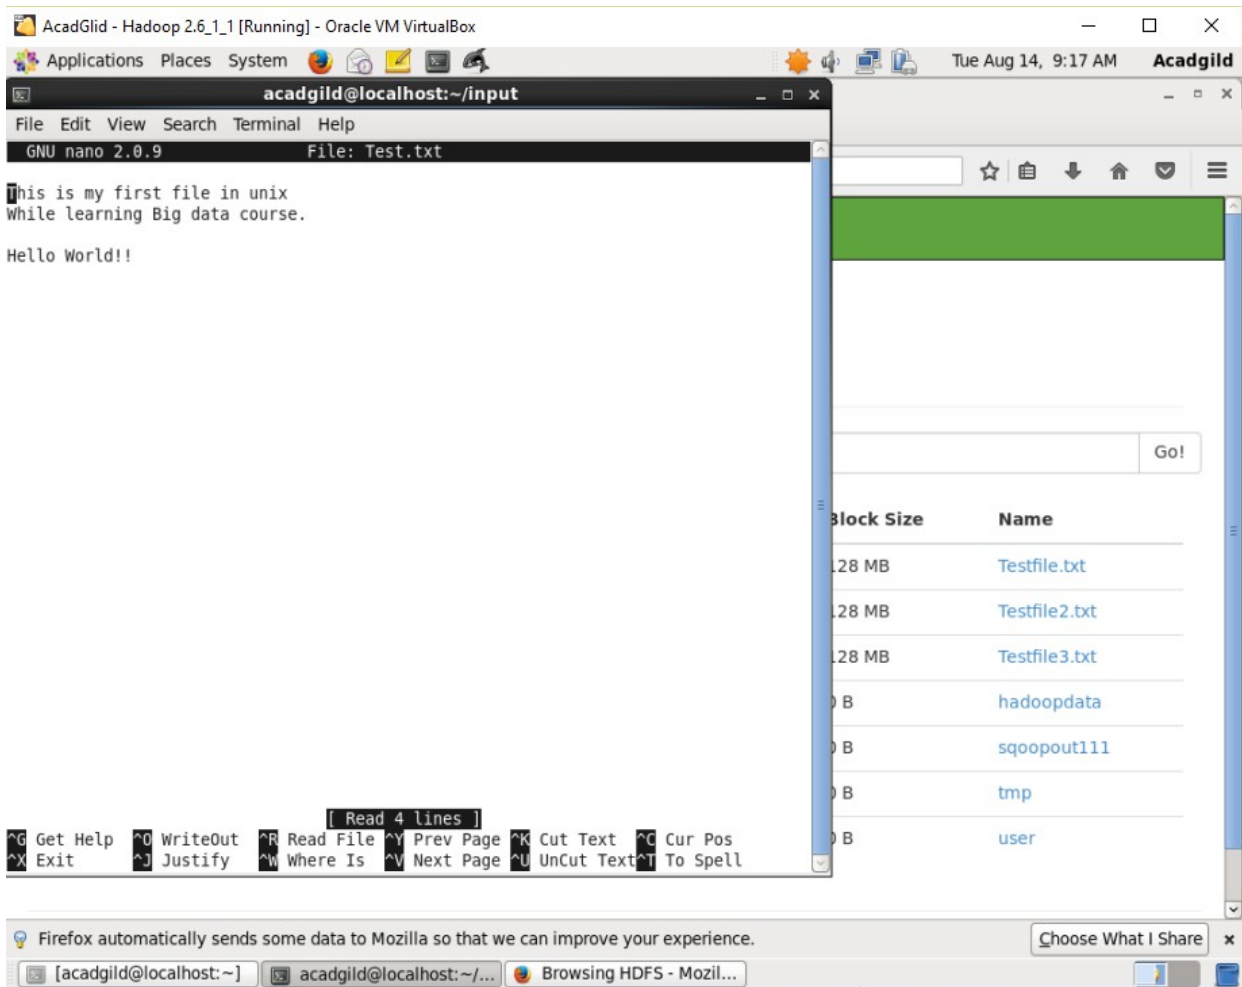
2. Run JPS will show started Java processors running post start up of hadoop processes.



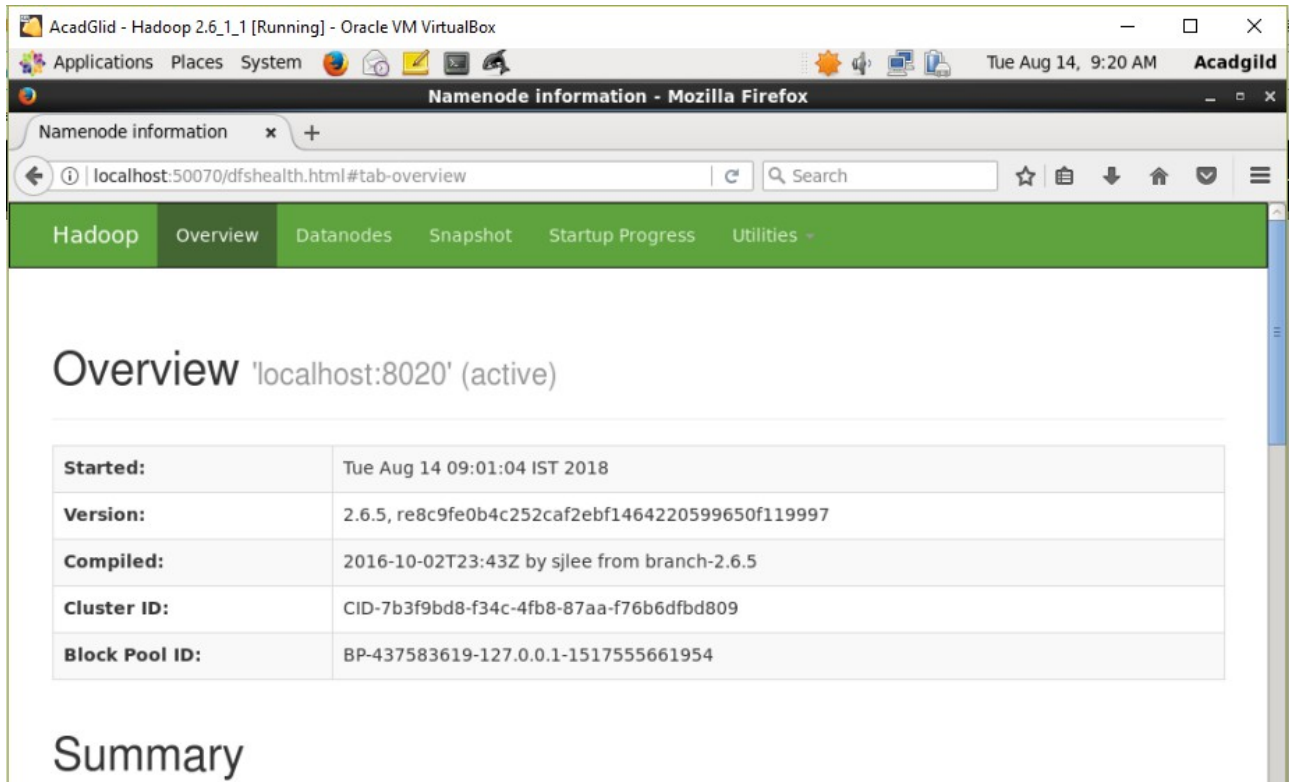
3. Run unix commands like pwd, ls -ls, ps, df etc..
For running some useful commands of unix necessary for basic operations like creating files, seeing current folder and directory list, processes running list, and disk file usage. So on unix local machine.



4. Use nano Test.txt to create file and cat out the contents
For creating files on local unix which can be uploaded to HDFS later on



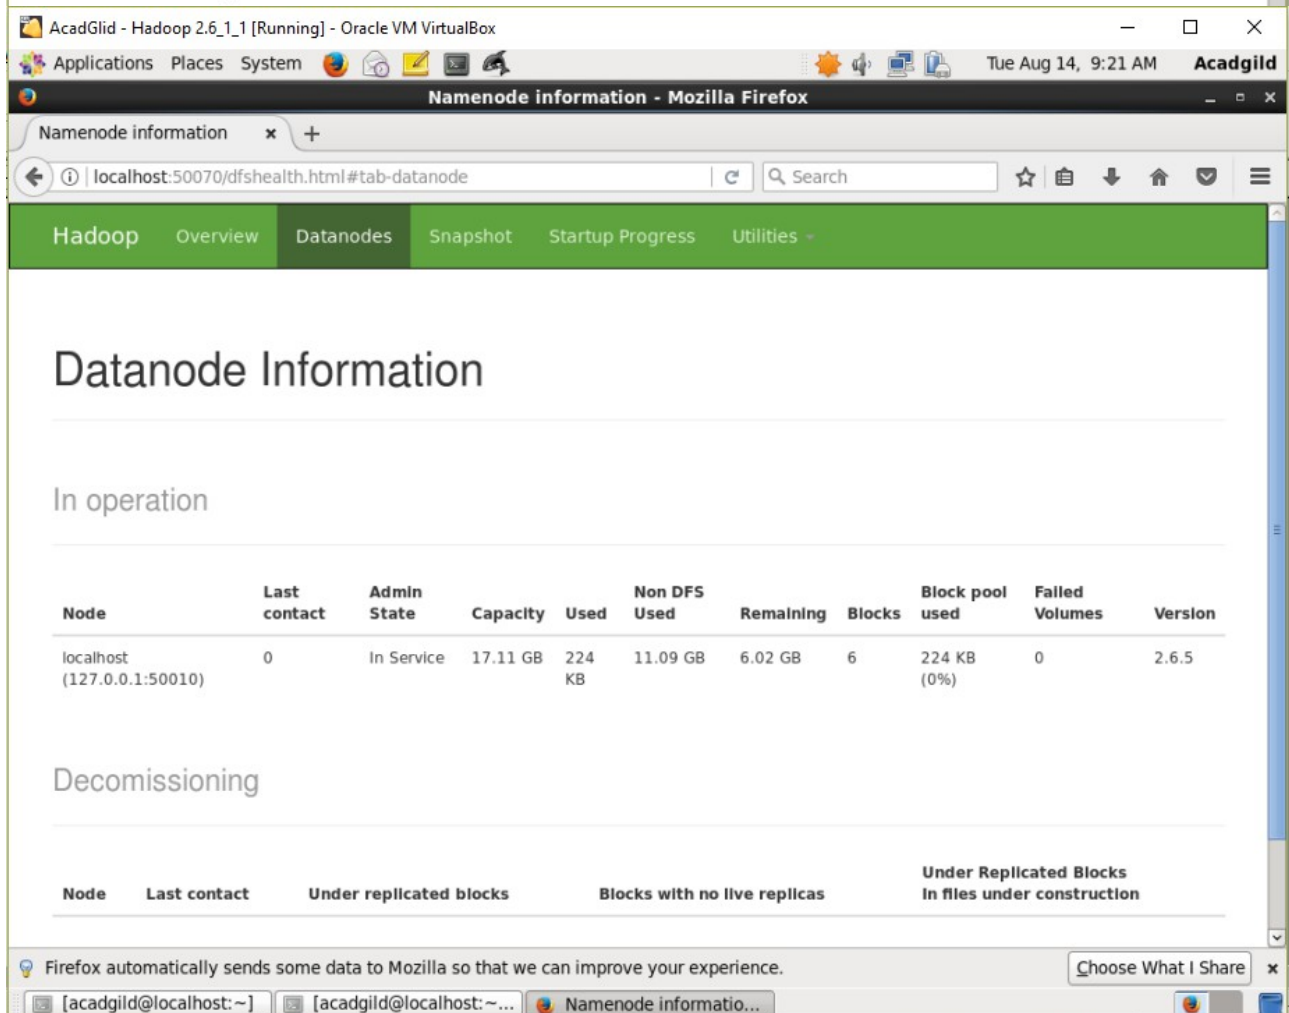
5. Rin <http://localhost:50070> and check out put
Can see status of Namenode, Data node and also see files in HDFS file system browsing like explorer. Also can see clicking on file status of file like replication factor and block size etc.
File saved in how many blocks and details of each block if it is stored in different data nodes.



The screenshot shows a web browser window titled "AcadGlid - Hadoop 2.6_1_1 [Running] - Oracle VM VirtualBox". The browser is Mozilla Firefox, and the address bar shows "localhost:50070/dfshealth.html#tab-overview". The page has a green navigation bar with tabs: "Hadoop", "Overview", "Datanodes", "Snapshot", "Startup Progress", and "Utilities". The "Overview" tab is selected. The main content area is titled "Overview 'localhost:8020' (active)". Below the title is a table with the following information:

Started:	Tue Aug 14 09:01:04 IST 2018
Version:	2.6.5, re8c9fe0b4c252caf2ebf1464220599650f119997
Compiled:	2016-10-02T23:43Z by sjlee from branch-2.6.5
Cluster ID:	CID-7b3f9bd8-f34c-4fb8-87aa-f76b6dfbd809
Block Pool ID:	BP-437583619-127.0.0.1-1517555661954

Below the table is a section titled "Summary".



The screenshot shows the same web browser window, but the address bar now shows "localhost:50070/dfshealth.html#tab-datanode". The "Datanodes" tab is selected in the navigation bar. The main content area is titled "Datanode Information". Below the title is a section titled "In operation". Below this is a table with the following information:

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
localhost (127.0.0.1:50010)	0	In Service	17.11 GB	224 KB	11.09 GB	6.02 GB	6	224 KB (0%)	0	2.6.5

Below the table is a section titled "Decommissioning". Below this is a table with the following information:

Node	Last contact	Under replicated blocks	Blocks with no live replicas	Under Replicated Blocks In files under construction
------	--------------	-------------------------	------------------------------	---

At the bottom of the browser window, there is a message: "Firefox automatically sends some data to Mozilla so that we can improve your experience." and a button "Choose What I Share".

AcadGild - Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox

Applications Places System Tue Aug 14, 9:24 AM Acadgild

Browsing HDFS - Mozilla Firefox

Browsing HDFS x +

localhost:50070/explorer.html#/

Hadoop Overview Datanodes Snapshot Startup Progress Utilities

Browse Dir

/

Permission	Name
-rW-r--r--	Testfile.txt
-rW-r--r--	Testfile2.txt
-rW-r--r--	Testfile3.txt
drwxr-xr-x	hadoopdata
drwxr-xr-x	sqoopout111
drwxrwx---	tmp
drwxr-xr-x	user

File information - Testfile3.txt

[Download](#)

Block information -- Block 0

Block ID: 1073741871
Block Pool ID: BP-437583619-127.0.0.1-1517555661954
Generation Stamp: 1047
Size: 18
Availability:

- localhost

Close

Firefox automatically sends some data to Mozilla so that we can improve your experience. Choose What I Share Close this message

[acadgild@localhost:~] [acadgild@localhost:~...] Browsing HDFS - Mozil...